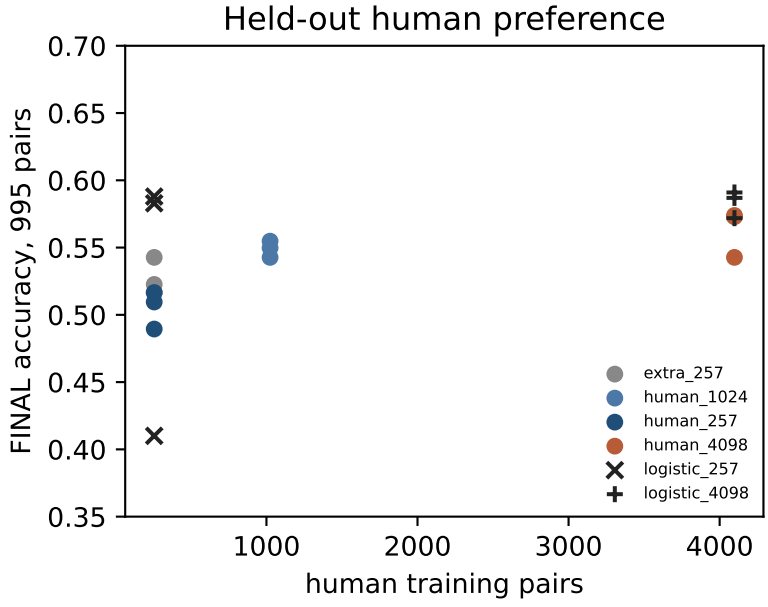
Held-out human preference
FINAL accuracy, 995 pairs
human training pairs
0.70
0.65
0.60
0.55
0.50
0.45
0.40
0.35
1000
2000
3000
4000
extra_257
human_1024
human_257
human_4098
logistic_257
logistic_4098

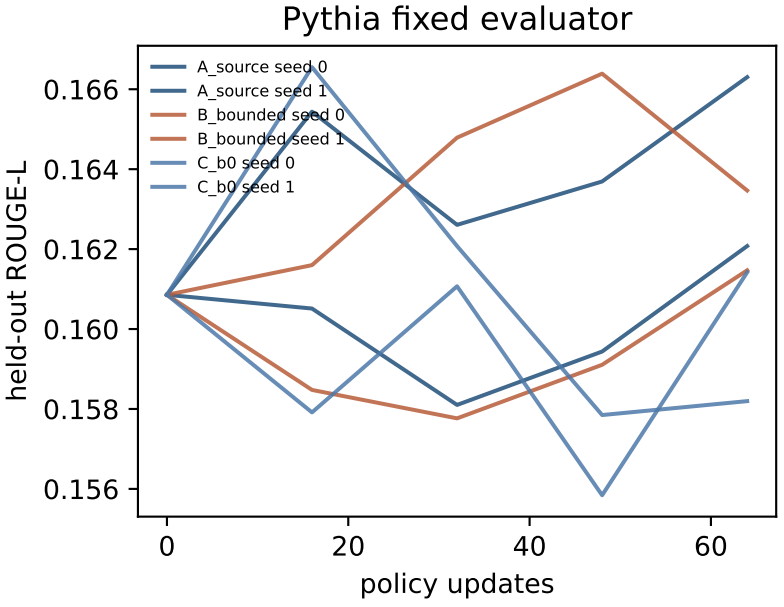
Pythia fixed evaluator
held-out ROUGE-L
policy updates
0.166
0.164
0.162
0.160
0.158
0.156
0
20
40
60
A_source seed 0
A_source seed 1
B_bounded seed 0
B_bounded seed 1
C_b0 seed 0
C_b0 seed 1